

Heuristic Search

Alice Gao

Lecture 3

Readings: RN 3.5 (esp. 3.5.2), PM 3.6

Outline

Learning Goals

Why Heuristic Search

LCFS, GBFS, and A*

- Lowest-Cost-First Search

- Greedy Best-First Search

- A* Search

Designing an Admissible Heuristic

Pruning the Search Space

Learning goals

By the end of the lecture, you should be able to

- ▶ Describe motivations for applying heuristic search algorithms.
- ▶ Trace the execution of and implement the Lowest-cost-first search, Greedy best-first search and A* search algorithm.
- ▶ Describe properties of the Lowest-cost-first, Greedy best-first and A* search algorithms.
- ▶ Design an admissible heuristic function for a search problem. Describe strategies for choosing among multiple heuristic functions.
- ▶ Describes strategies for pruning a search space.

Learning Goals

Why Heuristic Search

LCFS, GBFS, and A*

Designing an Admissible Heuristic

Pruning the Search Space

Why Heuristic Search?

How would ____ choose which one of the two states to expand?

- ▶ an uninformed search algorithm
- ▶ humans

5	3	
8	7	6
2	4	1

1	2	3
4	5	
7	8	6

Why Heuristic Search

An uninformed search algorithm

- ▶ considers every state to be the same.
- ▶ does not know which state is closer to the goal.
- ▶ may not find the optimal solution.

An heuristic search algorithm

- ▶ uses heuristics to estimate how close the state is to a goal.
- ▶ try to find the optimal solution.

Learning Goals

Why Heuristic Search

LCFS, GBFS, and A*

- Lowest-Cost-First Search

- Greedy Best-First Search

- A* Search

Designing an Admissible Heuristic

Pruning the Search Space

The Cost Function

Suppose that we are executing a search algorithm and we have added a path ending at n to the frontier.

$cost(n)$ is the actual cost of the path ending at n .

The Heuristic Function

Definition (search heuristic)

A **search heuristic** $h(n)$ is an estimate of the cost of the cheapest path from node n to a goal node.

In general, $h(n)$ can be arbitrary.

However, a good heuristic function has the following properties.

- ▶ problem-specific.
- ▶ non-negative.
- ▶ $h(n) = 0$ if n is a goal node.
- ▶ $h(n)$ must be easy to compute (without search).

LCFS, GBFS, and A*

- ▶ LCFS: remove the path with the lowest cost $cost(n)$.
- ▶ GBFS: remove the path with the lowest heuristic value $h(n)$.
- ▶ A*: remove the path with the lowest cost + heuristic value $cost(n) + h(n)$.

Learning Goals

Why Heuristic Search

LCFS, GBFS, and A*

- Lowest-Cost-First Search

- Greedy Best-First Search

- A* Search

Designing an Admissible Heuristic


Pruning the Search Space


Lowest-cost-first search

- ▶ Frontier is a priority queue ordered by $cost(n)$.
- ▶ Expand the path with the lowest $cost(n)$.

Trace LCFS on a search graph

If there is a tie, remove nodes from the frontier in alphabetical order.






LCFS

frontier: ~~S~~ ~~B~~ ~~C~~ ~~D~~ ~~E~~ ~~H~~ ~~H~~ ~~F~~ ~~F~~ ~~G~~

0 1 1 2 10 2 2 3 3 11 11



Properties of LCFS

- Complexity:

time and space complexities are both exponential.

- Complete: Guaranteed to find a solution if one exists?

Yes.

- Optimal: Guaranteed to find the optimal solution?

Yes.

mild technical conditions required:

- ① branching factor is finite. *constant.*
- ② the cost of every arc is bounded below by a positive \wedge

Learning Goals

Why Heuristic Search

LCFS, GBFS, and A*

Lowest-Cost-First Search

Greedy Best-First Search

A* Search

Designing an Admissible Heuristic


Pruning the Search Space


Greedy Best-First Search

- ▶ Frontier is a priority queue ordered by $h(n)$.
- ▶ Expand the node with the lowest $h(n)$.

Trace GBFS on a search graph

If there is a tie, remove nodes from the frontier in alphabetical order.






GBFS

frontier: ~~S~~ ~~B~~ ~~C~~ H C ~~D~~ E F ~~G~~

8	7	3	100	3	1	4	6	0
---	---	---	-----	---	---	---	---	---



Properties of GBFS

- ▶ Complexity:
space and time complexity are both exponential.
- ▶ Complete: Guaranteed to find a solution if one exists?
No.
- ▶ Optimal: Guaranteed to find the optimal solution?
No.

Learning Goals

Why Heuristic Search

LCFS, GBFS, and A*

Lowest-Cost-First Search


Greedy Best-First Search

A* Search

Designing an Admissible Heuristic

Pruning the Search Space


A* Search




- ▶ Frontier is a priority queue ordered by $f(n) = cost(n) + h(n)$.
- ▶ Expand the node with the lowest $f(n)$.


Trace A* search on a search graph

If there is a tie, remove nodes from the frontier in alphabetical order.





A^*
 frontier: ~~S~~ ~~B~~ ~~C~~ H ~~D~~ ~~E~~ H ~~F~~ ~~G~~
 8 8 4 102 5 11 6 103 9 17 11



Properties of A*

- Complexity:

space and time complexities are both exponential.

- Complete: Guaranteed to find a solution if one exists?

Yes, if the heuristic satisfies a mild condition.

- Optimal: Guaranteed to find the optimal solution?

Yes, if the heuristic satisfies a mild condition.

A* is Optimal

$h^*(n)$: cost of cheapest path from n to a goal.

$$0 \leq h(n) \leq h^*(n).$$

Definition (admissible heuristic)

A heuristic $h(n)$ is **admissible** if it is never an overestimate of the cost of the cheapest path from node n to a goal node.

Theorem (Optimality of A*)

If the heuristic $h(n)$ is admissible, the solution found by A* is optimal.

$h(n)$ – a lower bound on the actual cost
– an optimistic estimate of the actual cost.

A* is Optimally Efficient

Among all optimal algorithms that start from the same start node and use the same heuristic, A* expands the fewest nodes.

A could be unlucky w/ tie breaking.*

Define optimal efficiency as

- expanding the minimum number of nodes n for which $f(n) \neq f^*$ where f^* is the cost of the cheapest path.*

Learning Goals

Why Heuristic Search

Greedy Best-First Search

A* Search

Designing an Admissible Heuristic

Pruning the Search Space

Some Heuristic Functions for 8-Puzzle

- ▶ Manhattan Distance Heuristic: h_1
The sum of the Manhattan distances of the tiles from their goal positions
- ▶ Misplaced Tile Heuristic: h_2
The number of tiles that are NOT in their goal positions

Both heuristic functions are admissible.

Initial State (S) $h_2(S) = 7$

5 ¹	3 ²	³
8 ⁴	7 ⁵	6 ⁶
2 ⁷	4 ⁸	1

$h_1(S)$
 $= 4 + 3$
 $+ 1 + 2$
 $+ 2 + 0$
 $+ 2 + 2$
 $= 16$

Goal State

1	2	3
4	5	6
7	8	

Constructing an Admissible Heuristic

- ▶ Define a relaxed problem by simplifying or removing constraints on the original problem.
- ▶ Solve the relaxed problem without search.
- ▶ The cost of the optimal solution to the relaxed problem is an admissible heuristic for the original problem.

Constructing an Admissible Heuristic for 8-Puzzle

8-puzzle: A tile can move from square A to square B

- ▶ if A and B are adjacent, and
- ▶ B is empty.

Which heuristics can we derive from relaxed versions of this problem?

CQ: Constructing an Admissible Heuristic

CQ: Which heuristics can we derive from the following relaxed 8-puzzle problem?


A tile can move from square A to square B if A and B are adjacent. *(B may not be empty.)*

(A) The Manhattan distance heuristic

(B) The Misplaced tile heuristic

(C) Another heuristic not described above

$$4 + 3 + 1 + 2 + \dots$$



CQ: Constructing an Admissible Heuristic

CQ: Which heuristics can we derive from the following relaxed 8-puzzle problem? *(A and B may not be adjacent.)*


A tile can move from square A to square B. *(B may not be empty.)*

(A) The Manhattan distance heuristic


(B) The Misplaced tile heuristic

(C) Another heuristic not described above

$$\begin{aligned} &1+1+1+1 \\ &+ \dots + \end{aligned}$$



Which Heuristic is Better?



- ▶ We want a heuristic to be admissible.
- ▶ Prefer a heuristic that is very different for different states.
- ▶ Want a heuristic to have higher values (close to h^*).

Dominating Heuristic

Definition (dominating heuristic)

Given heuristics $h_1(n)$ and $h_2(n)$. $h_2(n)$ dominates $h_1(n)$ if

- ▶ $(\forall n (h_2(n) \geq h_1(n)))$.
- ▶ $(\exists n (h_2(n) > h_1(n)))$.


Theorem

If $h_2(n)$ dominates $h_1(n)$, A^ using h_2 will never expand more nodes than A^* using h_1 .*

CQ: Which Heuristic of 8-puzzle is Better?

CQ: Which of the two heuristics of the 8-puzzle is better?

- (A) The Manhattan distance heuristic dominates the Misplaced tile heuristic.
- (B) The Misplaced tile heuristic dominates the Manhattan distance heuristic.
- (C) I don't know....



Learning Goals

Why Heuristic Search

Greedy Best-First Search


A* Search

Designing an Admissible Heuristic

Pruning the Search Space


Cycle Pruning

- ▶ A cycle cannot be part of a least-cost path.
- ▶ Works well with depth-first search.
- ▶ The complexity of cycle pruning is ...



Multiple-Path Pruning

- ▶ If we have already found a path to a node, we can prune other paths to the same node.
- ▶ Subsumes a cycle check.
- ▶ Requires storing all nodes we have found paths to.
- ▶ What if a subsequent path to n is shorter than the first path found?




Lowest-cost-first search w/ multiple-path pruning

If we perform multiple-path pruning with lowest-cost-first search, is it possible for us to prune the optimal solution (least-cost path)?

(A) Yes, it is possible.

(B) No, it is not possible.

$$\text{cost}(P_1) > \text{cost}(P_2)$$



A* search w/ multiple-path pruning

Q: come up w/ a search graph s.t.


A* w/ multi-path pruning discards
the optimal solution.


If we perform multiple-path pruning with A* search,
is it possible for us to prune the optimal solution (least-cost path)?

(A) Yes, it is possible.

(B) No, it is not possible.

This example is
wrong. Check
the pages below
for correct
examples






explored: S, A, B, C, G

frontier: ^①S, ^②SA, ^③SB, ^④SAC, ^⑤SBC, ~~SBCG~~

7 9 31 21 21

path found: SBCG (21) optimal solution
SACG (31)




explored: S, A, B, G

frontier: ^①S, ^②SA, ^③SB, ^⑤SAG, ^④SBA

21 22 42 3

path found: SAG (42)
optimal solution: SBAG (24)



explored: S, C, D, A, B, G

frontier: $\textcircled{1}$ \$, $\textcircled{14}$ ~~SA~~, $\textcircled{2}$ ~~SG~~, $\textcircled{3}$ ~~SCD~~, $\textcircled{7}$ ~~SCDG~~, $\textcircled{5}$ ~~SAB~~, $\textcircled{6}$ ~~SABD~~

23 7 17 26 23 14

path found: SCDG (26)

optimal solution: SABDG (23)

Find Optimal Solution w/ Multiple-Path Pruning

What if a subsequent path to n is shorter than the first path found?

- ▶ Remove all paths from the frontier that use the longer path.
- ▶ Change the initial segment of the paths on the frontier to use the shorter path.
- ▶ Make sure that the least-cost path to a node is found first.

A* search w/ multiple-path pruning

How can we ensure that A* with multiple-path pruning is optimal?

- ▶ Ensure that we find the least-cost path to every node first.
- ▶ Admissible heuristic guarantees the above for a goal node, but not for other nodes. *→ consistent.*
- ▶ We need the heuristic to satisfy the monotone restriction:

for any arc $\langle m, n \rangle$, $\underline{h(m)} - \underline{h(n)} \leq \underline{\text{cost}(m, n)}$.

if n is a goal node, $h(m) \leq \text{cost}(m, n) \rightarrow$ admissible heuristic.

If the heuristic satisfies the monotone restriction,

A* search with multiple-path pruning is optimal.

$h(m) - h(n) =$ the heuristic estimate of the path cost from m to n .

Summary of Search Strategies

Strategy	Frontier Selection	Halts?	Space	Time
Depth-first	Last node added	No	Linear	Exp
Breadth-first	First node added	Yes	Exp	Exp
Lowest-cost-first	$\min \text{cost}(n)$	Yes	Exp	Exp
Greedy Best-first	$\min h(n)$	No	Exp	Exp
A*	$\min \text{cost}(n) + h(n)$	Yes	Exp	Exp

Revisiting the learning goals

By the end of the lecture, you should be able to

- ▶ Describe motivations for applying heuristic search algorithms.
- ▶ Trace the execution of and implement the Lowest-cost-first search, Greedy best-first search and A* search algorithm.
- ▶ Describe properties of the Lowest-cost-first, Greedy best-first and A* search algorithms.
- ▶ Design an admissible heuristic function for a search problem. Describe strategies for choosing among multiple heuristic functions.
- ▶ Describes strategies for pruning a search space.